

Geonomics: A Python package for building agent-based, spatially explicit, and arbitrarily complex landscape-genomic simulations

Drew Hart

February 22, 2019

1 Abstract

TO BE COMPLETED

2 Intro/Background

There is ever-growing interest in understanding and even predicting the genomic evolution of complex study systems on complex and changing landscapes. Such systems might include one or multiple populations or species that are not at equilibrium. These species might inhabit complex, multivariate, and even changing landscapes. And may be undergoing both neutral evolution and natural selection, often on multiple traits of variable genetic architecture. Landscape genomics studies the ways in which ecological and evolutionary processes playing out on real landscapes generate geographical patterns of genomic diversity, and the field frequently features analysis of data collected from study systems of such genomic and geospatial complexity. Study of such systems is crucial for improving our understanding of real-world systems [2, 20], and for better anticipating evolutionary responses to climate change [3] and other sources of environmental change in the Anthropocene.

But the complex genomics of such systems are beyond the reach of analytical population genetics, and their spatial complexity and multifaceted evolutionary dynamics make them intractable for coalescent simulation. This hinders not only our understanding of many empirical systems and our ability to unambiguously interpret analytical results, but also our ability to predict those systems' dynamics, and thus to manage them appropriately. Thus, as is increasingly the case in many fields, forward-time simulation is a crucial tool for dissecting the evolutionary dynamics of complex study systems in landscape genomics. However, the current suite of forward-time genomic simulators, however numerous, is still of limited use for such work. Most available software is limited, either genetically or geospatially, in the complexity it can model. Many programs can model systems of considerable genomic complexity (e.g. simuPOP [21], NEMO [6] and QuantiNemo [17]), yet incorporate no or only rudimentary spatial components. Various other programs are designed specifically for landscape-genetic simulations (e.g. CDPOP [14], CDmetaPOP [13], SimAdapt [23]), but are limited in their genomic complexity. For instance, many programs are incapable of modeling simultaneous selection on numerous multigenic traits. To our knowledge, SLiM [8, 16] is the only package currently capable of simulating scenarios that are both as genetically and as geospatially complex as those for which Geonomics is designed [11]. (Indeed, the generalizability and complexity of which SLiM is capable far exceeds that of Geonomics.)

We developed Geonomics to provide a scriptable framework for building complex landscape-genomic simulations with minimal effort. Geonomics models are parameterized by way of an informatively annotated parameters file that provides the user a straightforward means of building models of arbitrary complexity yet offers reasonable default settings and/or ‘off switches’ for model components that are not the focus the user’s interest. The complex landscape-genomic models for which Geonomics is designed would require a considerable amount of work to script in SLiM’s Eidos framework (see examples of spatialized selection in

SLiM 2’s recipe book; e.g. section 14.11, page 288; [7]), but can be built in Geonomics with as little as three lines of code. For example, a Geonomics user could build a model of evolution with natural selection on multiple multigenic traits, on a multivariate landscape undergoing spatially inhomogeneous environmental change for certain landscape layers, in a species moving realistically across that landscape, and then run this model an arbitrary number of times, collecting data at various points during each run — all of this by doing nothing more than templating a parameters file, making alterations, loading the file, and running the resulting model object.

What’s more, Geonomics is written and run entirely in Python, a broad-purpose and popular programming language that is already familiar to most researchers with exposure to bioinformatics. This of course makes Geonomics considerably slower than its brethren that are written in compiled languages such as C++ (e.g. SLiM). But run time is not expected to be a major constraint for the sorts of models for which Geonomics was built (see runtime analyses, below). And what Geonomics sacrifices in performance, it gains in flexibility, extensibility, and accessibility. In fact, the basic user needn’t even need to know how to write full Python scripts to build a Geonomics model; users can build their own models by recycling and tweaking a minimal amount of code (available within the documentation and at the Geonomics homepage). On the other hand, advanced users wishing to code their own extensions or customizations have broad opportunity to do so, because Geonomics is a Python package that is seamlessly integrated into the universe of Python functionality.

3 Model overview

3.1 Components

A Geonomics model consists of two core components: the species and the landscape.

3.1.1 Species

A species consists of an arbitrary number of individuals, which do not belong to discrete populations but instead are distributed in continuous space upon the landscape. A species is described by a wide variety of demographic and life-history parameters that determine its behavior in the model (e.g. intrinsic growth rate, mate-search radius, mean number of offspring per mating event, reproductive age and maximum age, and so on; for a detailed discussion and these and all other model parameters, see the documentation). Each species can also undergo any number of arbitrarily complex demographic changes during each model run, including both population-size changes (which can be exponential, cyclical, stochastic, or custom-defined) and changes to various demographic and life-history parameters.

Each individual in a species is described by an x,y location, a sex, an age (or stage), a genome (optionally), and a phenotype for any traits assigned to the species. (It is worth noting that species are collected into objects called communities, which for most models will consist only of a single species, but which provides a framework for the advanced user to code inter-species interactions for multi-species models; this is a functionality that we hope to build into future versions of the software.)

Each individual has a diploid genome consisting of a L diallelic loci. These loci can be treated as representing either a contiguous haplotype block or a set of discrete markers, depending on the map of recombination rates assigned to the genome. (For simplicity’s sake, we refer to these loci herein, and in the software generally, as ‘the genome’.) Genomes are initially assigned based on a species’ genomic architecture—an object containing parameters describing all simulated loci. These parameters include the starting allele frequencies and dominance values for all loci, the inter-locus recombination rates across the genome. Separate chromosomes can be modeled by providing a list of chromosome lengths, between which the recombination rate is set to 0.5, ensuring independent assortment.

A genomic architecture can also stipulate any number of mono- or multigenic traits for a species. Each trait is defined by a set of loci that underpin it, the effect sizes of those loci, and a selection coefficient (which can be heterogeneous or homogeneous in both space and time). Mutations, which are of three types:

neutral; deleterious, which universally decrease an individual's fitness); and trait-affecting, which influence an individual's phenotype, with the resulting fitness effect determined by the individual's local environment. All three mutation types are controlled by type-specific mutation rates (additional parameters within the genomic architecture, any or all of which can be set to zero). (To simulate complex, specific genomic architectures, users can feed into Geonomics a CSV-formatted file defining the architecture locus by locus. For details, see the documentation.)

3.1.2 Landscape

Aside from the species, the other core component in a Geonomics model is the landscape. A landscape is a stack of an arbitrary number of layers (i.e. variables), each represented by a raster of normalized values ($0 \leq \text{value} \leq 1$). Each layer can be programmed to serve as the basis for any of a number of model components: 1. the raster of cell-wise carrying-capacities controlling the population density of a species; 2. the resistance surface controlling the realistic movement of individuals and/or dispersal of offspring across the landscape; 3. the selective force acting on one or more traits of one or more species.

Each layer of a landscape can be programmed to undergo any number of arbitrarily complex environmental change events during each model run. The changes they produce will in turn affect the dynamics of any species on the landscape for which that layer plays a role in its population dynamics, movement or dispersal, and/or natural selection. Each event is defined by the timesteps over which it unfolds, and either the terminal raster of the event (with intermediate rasters being linearly interpolated) or a directory containing the stepwies time-series of rasters pertaining to the event. The latter option makes it extremely easy to simulate evolution on real-world rasters undergoing real, non-linear and spatially heterogeneous environmental change.

3.2 Operations

A Geonomics model can be parameterized to run for an arbitrary number of runs, with each run creating a separate subdirectory of output. At the start of each run, the must be burned in. This is accomplished by running the model (without the genomic or selective components) until a series of statistical tests (a time-lagged t-test of mean population size; an augmented Dickey-Fuller test of population size; STILL NEED TO ADD THE SPATIAL TEST) determines that each species' population size and spatial distribution has reached dynamic equilibrium. Then, if genomes are being used, each individual has a genome randomly drawn (according to its genomic architecture) and assigned, such that the main phase of each model run begins in the absence of any population structure. At this point, the 'main' phase of the run can run for any number of timesteps. Each timestep is composed of a series of actions, some requisite, some optional (see Figure 1 for details):

1. age/stage incrementation (requisite);
2. movement (optional);
3. mate-finding and mating (requisite);
4. gamete production (optional, because use of genomes is optional);
5. offspring dispersal (requisite);
6. mortality (due to the combination of density-dependence [requisite] and natural selection [optional]);
7. demographic change events (optional);
8. landscape-change events (optional);
9. recording of data and statistics (optional).

Each individual's age/stage increments at each timestep. A number of parameters can be set so as to modify a model's behavior on the basis of this attribute (including the minimum age of reproduction and the maximum age of a species).

Movement takes place in continuous space, such that individuals are not arbitrarily restricted to grid cells. Individuals' distances and directions are drawn separately, then composed into movement vectors. Distances are Wald-distributed (and the distributional parameters, as with nearly all distributions used in the model, can be set by the user within a Geonomics parameters file). Directions can either be drawn from a uniform distribution on the unit circle, resulting in isotropic movement (the default behavior); or they can be drawn from a 'movement surface' — an array of uni- or multimodal Von Mises distributions, derived from a landscape layer that serves as a resistance surface. The latter option is, to our knowledge, a novel approach to simulating movement, which generates realistic, anisotropic movement across an environment of heterogenous habitat quality. (For details, see Figures 2 and 3).

Mating pairs are chosen from among all pairs of individuals within the species' mate-search radius (based on eligibility by age and sex, and decided by a Bernoulli draw with probability equal to the species' intrinsic birth rate). For each mating pair a number of offspring is chosen (from a Poisson distribution with lambda equal to the species' mean number of offspring, unless the user fixes the number of offspring lambda). Each parent produces one gamete for each of its offspring, by recombination (at the inter-locus rates defined by the species' genomic architecture) and Mendelian segregation. Offspring individuals are created and then dispersed to a new location (where, as with movement, the directions of their dispersal vectors can be drawn either isotropically or anisotropically, the latter using a 'dispersal surface').

Mating is followed by mortality. Deaths are drawn binomially, based on individual-wise death probabilities, which are calculated as a combination of a combination of the probability of death by density-dependence (from a spatialized logistic-growth model) and the probability of death by natural selection (on any number of traits simultaneously). This is calculated as:

$$P(d_i) = 1 - (1 - P(d_{x,y})) \prod_{p=1}^m \omega_{i,p} \quad (1)$$

where $P(d_{x,y})$ is individual i 's probability of death by density-dependence and $\omega_{i,p}$ is individual i 's fitness for trait p (and only factors in if natural selection is being used). The probability of density-dependent death at location x, y is calculated as:

$$P(d_{x,y}) = E[N_{d;x,y}] / N_{x,y} = \frac{E[N_{b;x,y}] - \frac{dN_{x,y}}{dt}}{N_{x,y}} \quad (2)$$

where $E[N_{d;x,y}]$ is the expected number of deaths at the individual's x, y location on the landscape; $N_{x,y}$ is the population density (expressed as individuals per cell) at the location; $E[N_{b;x,y}]$ is the expected number of births at the location; and $\frac{dN_{x,y}}{dt}$ is the logistic population growth rate at the location. Individual i 's fitness for trait p is calculated as:

$$\omega_{i,p} = 1 - \phi_{p;x,y} (|e_{p;x,y} - z_{i;p}|)^{\gamma_p} \quad (3)$$

where $\phi_{p;x,y}$ is the selection coefficient on trait p at the individual's location; $e_{p;x,y}$ is the environmental value of the layer that serves as the selective force for trait p ; γ_p defines the curvature of the fitness function for trait p , and $z_{i;p}$ is individual i 's phenotype for trait p , which is calculated from the additive effects of the individual's genotypes at all influencing loci (Geonomics does not model epistasis) as:

$$z_{i;p} = g_0 + \sum_{l=0}^n \alpha_{p,l} g_{i,l} \quad (4)$$

where n is the number of loci, α is a locus' effect size, g is the individual's genotype, and *baseline_genotype* is 0 for monogenic traits, 0.5 for polygenic traits.

Geonomics is designed as an object-oriented scripting framework with both basic and advanced use-modes. The basic mode simply requires users to call a first command to create a template parameters file

(which they must then edit as desired), a second command to create a model object from that parameters file, and a third command to then run the desired number of runs for that model (see Figure 1). The advanced mode allows users to make modifications to the components of their models, or to collect custom data from their model, by calling additional functions before a model is run, between the timesteps of a run, or after a run is complete. This can be done using both built-in Geonomics functions and homespun Python code. (This is what makes Geonomics so extensible — because it is a Python package, users can call on the full spectrum of Python functionality to design custom code that can interact with Geonomics objects.)

Landscape and demographic change events unfold over some portion of the total time of a run. The changes are made incrementally, with each incremental change being made during each of the timesteps in the series of timesteps defined *a priori* by the user (in the parameters file). Likewise, statistics are calculated and data are collected at each of the set of timesteps defined *a priori* by the user.

4 Validations tests

5 Validations tests OLD

We have run a series of tests to statistically and heuristically validate the full range of functionality available in Geonomics. In this section, we briefly review the reasoning and parameterization for each test, then present the results. (We discuss only parameters of key importance to each scenario. Other parameters are set to reasonable values, or left at defaults. For full details, see the parameters file for a given test.)

All results are as expected, demonstrating a robust ability to reproduce population-genetic and population-genomic findings. Some results show minor deviations from theory. These artefacts are primarily a resulting of using a simulation framework designed for complex, spatially explicit models to approximate much simpler and, in some cases, aspatial or spatially implicit models. They artefacts are discussed where applicable. (The code for all tests is available in the ‘./tests/validation’ directory of the package.)

5.1 Wright-Fisher test: genetic drift

The Wright-Fisher model of genetic drift models a fixed-size haploid population that turns over completely at each timestep (i.e. generation). The population can have any number of independent, biallelic genetic loci. For each locus, each generation’s allele frequency is chosen as a binomial random variable, with the number of trials equal to the population size and the probability of success (i.e. of drawing the ‘1’ allele) equal to the previous generation’s ‘1’-allele frequency. The mean persistence time for an allele (i.e. the expected number of generations for which a locus remains segregating) is:

$$\bar{t}(p) = -4N[(1-p)\ln(1-p) + p\ln(p)] \quad (5)$$

where $2N$ is the number of alleles in the population (such that N can represent the diploid population size) and p is the frequency of either allele at the locus [4, 9, 26].

Clearly, the Wright-Fisher model is much simpler than the sorts of models for which Geonomics is designed (as are all of the following validations tests)—it is aspatial, panmictic, features fixed population sizes, models only neutral loci, and so forth. Thus, we parameterized Geonomics so as to approximate the model as closely as possible. To emulate aspatiality and panmixia, we used a population on a homogeneous landscape, with isotropic movement, and with movement and dispersal distributions and a mating radius that broadly encompass the diagonal width of the landscape. To enforce complete generational turnover, we set the maximum age parameter to 1 timestep. While Geonomics does not maintain constant population size, we maintained the carrying-capacity raster at a constant, uniform value, thus maintaining a stationary mean population size. We simulated 100, independent neutral loci (by setting all interlocus recombination rates to 0.5), with starting ‘1’-allele frequencies of 0.5 (although the actual starting frequencies vary slightly around this value because of sampling error when all individuals’ genotypes are drawn binomially).

We ran the Wright-Fisher approximation test for three values of the carrying-capacity raster (i.e. three values of ‘K_factor’), hence for three mean population sizes. For each mean population size (calculated as

the harmonic mean, to account for stochastic fluctuations around the carrying capacity), we compared mean persistence time to that expected by theory, according to equation 5 in the previous paragraph. As can be seen in Figures 4 and 5, the results are a close match to theory.

5.2 Bottleneck test: population dynamics

Because drift is a stronger evolutionary force in smaller populations, drift accelerates in shrinking populations. If a population undergoes a bottleneck event, the overall effect of drift on the population during that time is expected to be larger than a constant-size population of equivalent starting size would experience during that time. Thus, mean fixation time should decrease in a bottlenecked population relative to a constant-size comparison population.

As with the Wright-Fisher model, we used a homogeneous landscape with broad distributions for movement and dispersal and with a mating radius that encompasses the full landscape to emulate aspatiality and panmixia. To simulate a bottleneck event, we created a custom change event in which the population’s carrying-capacity raster is reduced to 30% of its initial value for 50 timesteps (from the 200th to 250th), then returned to its initial value for the remainder of the simulation (through the 2500th timestep).

Figure 6 shows a clear signal of drift acceleration during the bottleneck event.

5.3 Stepping-stone test: population subdivision and genetic differentiation

The stepping-stone model, or one-dimensional island model, is a spatially implicit model. It models a series of subpopulations, arranged along a straight line, with migration between all neighboring pairs. As a combined result of divergence by drift and homogenization by effective migration, subpopulations are expected to reach a stationary level of genetic differentiation—migration-drift equilibrium. Theory provides the expected pairwise genetic differentiation between a pair of subpopulations at equilibrium as:

$$F_{ST} = \frac{1}{1 + 4Nm} \quad (6)$$

where N is the population size and m is the per-generation migration rate, such that Nm can be interpreted as the per-generation number of migrant individuals [9].

To approximate the stepping-stone model, we created a Landscape Layer with a diagonal of six equally spaced islands (1.0-valued cells) embedded in a ‘sea’ of 0.0-valued cells. We used this layer as the carrying-capacity raster (Figure 9, left). We set the mating radius to encompass an individual’s current island, but no neighboring islands. We parameterized dispersal to be very local to parents’ midpoints, and parameterized movement distance to be strongly right-skewed, such that the long-distance movement events leading to migration are uncommon. We ran the simulation for 5000 timesteps.

Because Geonomics does not model discrete populations, it does not stipulate migration rates between discrete locations on the Landscape. Thus we manually tracked the number of migration events during each timestep, for all possible directional migration events (i.e. for all permutations of island pairs), then used that data to calculate all mean migration rates. With those values, we solved equation 6, then compared the resulting F_{ST} expectations to the observed values (calculated from the simulated data using two common methods; see Figure 9 for details).

Results demonstrate that the model has approached migration-drift equilibrium (Figure 7), and that all island populations were at dynamic equilibria around the same mean size, as expected by theory (Figure 8). Estimated migration rates and F_{ST} values qualitatively match theoretical expectations: mean migration rate drops off precipitously at greater than one island’s distance apart, and genetic differentiation increases to approximate saturation. Values of F_{ST} consistently undershoot the values expected based on estimated migration rates, however, because subpopulations have yet to approach fixation at most loci (which is the expectation implied by expected F_{ST} values close to 1).

5.4 Contrasting-habitat test: adaptive divergence

In a population divided between two opposite selective environments, if there is standing genetic variation for a biallelic locus controlling the trait responding to those environments, then theory predicts that the two subpopulations will diverge at that locus as each moves toward its respective adaptive peak. The rate at which divergence should occur depends on the relative strengths of two opposing evolutionary forces: natural selection, which causes divergence, and gene flow from migration, which causes homogenization. The rate of allele frequency change in either subpopulation at timestep t is expressed as:

$$\delta q = \frac{-spq[q + h(p - q)]}{1 - sq(2hp + q)} + m_i q^* - m_o q \quad (7)$$

where q and p are the frequencies of the deleterious and beneficial alleles in the subpopulation, s is the selection coefficient against the homozygous recessive phenotype, h is the degree of dominance of the recessive allele, m_i and m_o are the migration rates into and out of the subpopulation being analyzed, and q^* is the frequency of the recessive allele in the alternative subpopulation [9].

This model, much like the stepping-stone model, is spatially implicit. To approximate this, we created a landscape with two layers. The first was divided into two equal-sized halves, one valued at 0.0, the other at 1.0; this layer was used as the layer driving natural selection. The second was valued uniformly at 1.0; this was used as the carrying-capacity raster (thus setting uniform population density across the Landscape and determining, in sum, the overall carrying capacity of the landscape). We created one monogenic trait whose position was randomly chosen within a genomic architecture of 100 independent (i.e. unlinked) loci. We ran the model for 1000 timesteps for each of three values of the parameter phi (identical to s in equation 7): 0.1, 0.05, and 0.01. Given that Geonomics does not employ express migration rates, we tracked the number of migration events (between the two contrasting halves of the first layer) during each timestep, then used that data to solve equation 7.

Results depict clear local adaptation to each of the two halves of the landscape, with opposite-phenotype bleedover and heterozygote births occurring along the border between the two habitats (Figure 10 right). Allele trajectories in each half of the environment follow qualitatively the increasing and saturating trajectories expected by theory, but reach consistently more extreme allele frequencies than expected (Figure 10, left).

5.5 Cline test: local adaptation

In a clinal model, a population adapts locally across an environmental gradient, which is characterized by the extremes of its environmental values and its steepness (i.e. the instantaneous rate of environmental change along it). Local adaptation across this gradient will generate a cline, i.e. a geographic gradient in allele frequency (though natural selection is not the only way a cline could be produced). The clinal pattern is only expected for loci that underlie the trait undergoing selection along the cline (and loci in linkage). Unlinked loci have no long-term clinal expectation (though they could initially be swept along with the selective locus adaptation, and any number could continue to show spurious concordant clinal variation). To detect climally adapted loci, we can fit cline curves to the allele-frequency variation across the environmental gradient for all loci, with the expectation that the clines fit to adaptive loci will mirror the gradient. Numerous equations have been used to fit clines, but one of the most common is the sigmoidal \tanh function:

$$p_x = \frac{1}{2}(1 + \tanh[\frac{2(x - c)}{w}]) \quad (8)$$

where p is the frequency of the reference allele at position x along the cline, c is the centerpoint of the cline (such that $p_{x=c} = 0.5$), and w is the ‘width’, which is defined as $w = \frac{1}{slope}$ at centerpoint c [22].

To implement the cline model in Geonomics, we created a landscape with two layers. The first layer was an environmental layer—a symmetrical, non-linear gradient between 0-valued and 1-valued halves (see raster in both halves of Figure 11). The second was a uniformly valued habitat-quality layer, used to set a uniform population density and thus determine the global carrying capacity. We created a monogenic trait

whose locus was randomly placed within a genomic architecture of 100 independent loci. The trait had a *phi* (i.e. s) of 0.01, with the gradient layer serving as its selective force.

We ran the cline model for 2500 timesteps, then used a numerical optimization function (in Python’s `scipy` package [12]) to fit *tanh* clines to all loci. We plotted all fitted clines on top of the first landscape layer, with the cline for the one selective locus highlighted. The selective locus consistently and clearly stands out as the only locus with a cline matching the expectation (i.e. mirroring the environmental gradient; Figure 11, left), and results show an obviously locally adapted population, with a zone of hybridization and phenotypic spillover surrounding the cline’s center (Figure 11, right). Furthermore, for a family of regression models of environmental value on genotype for all loci, after Bonferroni correction for multiple testing, the selective locus is consistently significant and the most significant (though other loci produce false-positive results, albeit with considerably larger p-values). INCLUDE AND SHOW THESE REGRESSION RESULTS?

5.6 Selective sweep test: genetic hitchhiking

Genomic context and linkage add important complication to models of molecular evolution. The most basic model of selection with linkage is that of the selective sweep: a beneficial mutation occurs in a population, falling on a random genomic background, then rises in frequency because of its selective advantage until it becomes fixed, pulling up the frequency of the surrounding haplotype block in the process. But as the haplotype block increases in frequency it is nevertheless subject to recombination, which gradually erodes it symmetrically around the beneficial mutation. Thus the selective-sweep model predicts that once a beneficial mutation occurs —as long as it is not lost early on by chance— it and the haplotype block around it will rise in frequency, the mutation will eventually fix, potentially with some core block around it, but the rest of the block will erode fade over time. The haplotype block should be clearly visible in genomic data, where it will manifest as a genomic region of reduced diversity and heterozygosity, centered on the mutation.

To implement the selective sweep model in Geonomics, we again created a model approximating an aspatial, panmictic population (see Wright-Fisher test for details). We created a single, monogenic trait with a *phi* (i.e. s) of 0.15. The trait’s locus was manually set to position 50, such that it was at the center of the the 101-locus genomic architecture. We manually set the starting ‘1’-allele frequency at this locus to 0.0, but set the trait to selected upon by the landscape’s first and only layer (a uniformly 1-valued raster), such that all individuals began the model equally ill-fit (i.e. with a fitness value of $1 - \phi = 0.85$). After burning the model in, we iteratively chose a random individual, introduced a ‘1’-mutation in their genome at locus 50, ran the model for 50 timesteps, and checked whether the ‘1’ allele had reached a frequency greater than 0.05 by that time. We iterated until that check was passed, at which point we declared the mutant allele ‘established’ and continued to run the model until 2500 timesteps after the mutant reached fixation. At three timepoints during that model we calculated and recorded genome-wide nucleotide diversity using a sliding-window approach.

Geonomics realistically emulated the behavior of a selective sweep. The adaptive phenotype (the ‘1’—‘1’ genotype, plotted as white on a white environmental background; in Figure 12, top row) clearly emerged in a region surrounding the mutation’s origin, then spread rapidly throughout the population. The population’s mean fitness increased quickly from 0.85 (the universal fitness value before the mutation was introduced) to 1.00 (the universal fitness value after the sweep was complete; Figure 12, bottom row, leftmost plot). The linkage block around the sweeping locus was clearly visible as a region of depressed nucleotide diversity, which became more pronounced as the sweep went to completion, then slowly eroded as a result of recombination of the mutant haplotype’s alleles onto non-mutant backgrounds (Figure 12, bottom row, second, third, and fourth plots from the left).

5.7 PCA test: isolation by resistance

Many real-world populations inhabit landscapes with complex patterns of heterogeneous habitat suitability. The probability of an individual moving across each part of these landscapes is a function of that part’s habitat. Ecologists often use resistance surfaces (or their reciprocal, resistance surfaces) to describe movement across such landscapes. Geonomics’ movement surfaces and dispersal surfaces both model such movement

(for individuals' timestep-to-timestep movement and for the dispersal of new offspring, respectively). In a population evolving on such a landscape, gene flow between any two locations on the landscape is expected to be an inverse function of the resistance-distance between the locations. As a result of this, a pattern of isolation by resistance (IBR) is expected to develop: pairwise genetic distances between different populations, or different regions, should be positively correlated with pairwise resistance distances.

To test this, we constructed a Geonomics model of a single species evolving neutrally for 1000 timesteps (i.e. 1000 rounds of mating) on a randomly generated complex landscape layer, with that layer serving as the basis for the species' movement surface. We ran a genetic Principal Components Analysis (PCA) on the full species' simulated genomes, both after the burn-in (i.e. when genomes had just been randomly drawn and assigned to all individuals) and after the model had run. For each PCA, we extracted the first three principal components (PCs), scaled them to $0 \leq \text{value} \leq$, then used the resulting numerals to determine a the red, green, and blue (RGB) values for the color of each individual. We used those colors to color each of the individuals in a plot of the full species (using the `mod.plot()` command).

The results (in Figure 13) show a clear lack of spatial structure at the outset; because genomes were randomly drawn and assigned, the population possessed no spatial structure. But more importantly, the results demonstrate significant spatial structure, which maps onto the movement surface's landscape layer exactly as one would expect. Neutral evolution with realistic movement across this landscape generated a clear and strong signal of IBR.

5.8 simultaneous-selection test: selection on multiple traits

FIGURE OUT WHY THESE RESULTS DON'T LOOK RIGHT One of the powerful features of Geonomics is that it can simulate on numerous traits simultaneously, each trait being selected upon by a distinct, and potentially spatially differentially distributed selective force. When a population is undergoing selection of this nature, the evolutionary outcome should be a function of the genomic architecture of the traits (i.e. how many loci underlie them, and whether or not their loci are linked), and of the correlation of the two traits' selective forces in space.

We simulated two scenarios for 1000 timesteps each. Both scenarios are situated on the same landscape—a landscape with two distinct and uncorrelated selective-force layers: one a symmetric, horizontal environmental gradient from 0 to 1, the other a vertical gradient. In both scenarios, the simulated species has a 20-length genome, with 10 distinct loci underpinning each of the two traits in the simulation (i.e. there is no pleiotropy, hence no overlap in loci between the two traits). However, all loci are independent in the first scenario (recombination rate = 0.5), whereas all loci are linked in the second scenario (recombination rate = 0.05).

Theory suggests that local adaptation should proceed successfully in the first scenario, but should be hindered in the second scenario because of genomic conflict resulting from opposing selection on closely linked loci. Heuristically, our results corroborate these expectations, with scenario 1 showing a clear signal of local adaptation (Figure 14) whereas scenario 2 shows a much messier and more complicated spatial distribution of phenotypes for each of the two traits (Figure 15).

6 Example use-case: Polygenic adaptation to climate change in the Yosemite region

- good for now, but reread and edit anyhow

7 Example use-case: Polygenic adaptation to climate change in the Yosemite region OLD

Genomics' use and value is best demonstrated by walking through a worked example. In this section, we explain a complex evolutionary scenario for which simulated data is desirable, explain the steps by which we

have simulated the scenario in Geonomics, and then present the results. (The source code for this example is available within the `./example/yosemite` directory of the Geonomics repository.)

There is growing interest in the evolutionary implications of climate change. Much of this interest focuses on species that are locally adapted along some environmental gradient that is expected to shift under climate change. Such environmental shifts can be non-homogeneous spatially (e.g. cooler regions at higher altitudes appear to be warming more quickly than warmer, low-altitude regions). Of particular interest the ability for such species to adapt to changing conditions rather than simply shifting their ranges and/or being driven locally extinct.

We simulate an example of such a study system: the adaptive response of a hypothetical, continuously distributed, locally adapted species to spatially heterogeneous climate change in the Yosemite region. To do so, we began by downloading a raster dataset of 30-year temperature normals for the Yosemite region (from Cal-Adapt; www.caladapt.org), cropped it to a 90-cell by 90-cell window surrounding the Yosemite valley, and saved the cropped file. We then wrote a couple short Python functions to produce from that raster the full set of rasters needed for the simulation (including current and projected future temperatures, and current and project future habitat suitability).

Temperature is the environmental variable driving natural selection. In order to model adaptive responses to change in this variable, we needed a raster of future temperatures, to feed into a landscape-change event as the endpoint raster. We could have downloaded a series of future-projection rasters and loaded them in as the stepwise changes for the environmental-change event, because Geonomics is designed to accept a directory of such files. However, to keep things simple we just processed the 30-year normals raster with a simple heuristic function. That function adds to the 30-year normals raster a fixed, region-wide temperature increase (2 degrees) plus an additional, elevation-dependent fraction of that increase (where that fraction varies from 0.0 at the hottest cell to 1.0 at the coldest cell). This function emulates the fact that warming due to climate change is quickest at higher elevations. We saved the resulting future-temperature raster to a new raster file.

We also needed habitat-suitability rasters for the model. These must reflect the fact that species tend to find highest-quality habitat, and thus exist at their highest densities, in the core of their range, and that edge habitat becomes increasingly marginal. Thus we wrote a function that accepts the temperature raster as input and produces from it a second habitat-suitability raster (where cells with temperature values between 7 and 11 degrees Celsius are assigned a 1.0 quality value, and cells outside that range are assigned suitability values that linearly decrease to 0.0 in the hottest and coldest cells in the raster). We fed both the 30-year normals raster and the future temperature raster through this function, thereby generating both current and future habitat-suitability rasters, which we saved to new raster files.

With these rasters prepared, we then created a Geonomics parameters file. We needed to parameterize a model with 2 ‘file’-type Layers (i.e. layers read in from files; see the Geonomics documentation for details), both of which would undergo landscape-change. We also needed to model to contain 1 Species, with movement, a movement-surface, and genomes with 1 trait. And we needed the model to carry out data-collection. The Geonomics command we ran to generate a template Geonomics parameters file with those components was:

```
>>> gnx.make_parameters_file(filepath='yosemite_params.py',
layers=[{'type': 'file', 'change': True},
{'type': 'file', 'change': True}],
species=[{'movement': True, 'movement_surface': True,
'genomes': True, 'n_traits': 1}],
data=True)
```

After running this command, we opened the resulting, auto-generated parameters file (`yosemite_params.py`) and edited the parameters as needed. In the Layer ‘init’ and ‘change’ parameters sections we replaced the placeholder filenames with the names of the four temperature-derived rasters whose creation is explained above. We also parameterized the landscape-change events: both Layers would change from their original rasters to their final rasters in 20 stepwise changes over 1000 timesteps (i.e. 1000 years), starting on timestep 499 and finishing on timestep 1499. We stated that the habitat-quality layer would serve as the basis for the

movement-surface. And we parameterized the model to include no mutation. We stipulated the selection coefficient (*phi*) on the trait (0.05), and the number of loci underlying it (100 randomly-selected loci). We left all other parameters at default values.

Finally, we wrote a short script to:

1. create a model from the parameters ('mod = gnx.make_model(''./yosemite_params.py'')');
2. manually burn in and run the model (using the `mod.walk` function to run the model in either 'burn' or 'main' mode for the desired number of timesteps;
3. plot the model at various timesteps, using `mod.plot_<>` functions or manual `matplotlib` code.

The model generated a clear and realistic pattern of polygenic adaptation to the elevation-based temperature gradient in the Yosemite region, and that gradient of local adpatation exhibited a pronounced upslope shift in response to the period of climate change. These results are visible both heuristically (from the individuals' phenotypes plotted across the three discrete timesteps' columns in Figure 16) and analytically (from the rasters of neighborhood-meaned phenotype plotted across Figure 16, row 3).

8 Caveats and considerations

- runtime analyses (basic), and explicitly point out the key concerns/parameters that users should think about for runtime and memory usage constraints (including, e.g., a layer that serves as basis for movement-surface and that changes a lot) - need to discuss the weird approach to mutation?

There are two unconventional approximations that Geonomics uses, in order to make such complex models tractable in Python within a reasonable amount of compute-time. The first is the approximation used to model recombination along a simulated genome. Precise modeling of recombination between all loci when each gamete is produced would require an extremely large number of Bernoulli draws—roughly $2LbT$, where 2 represents diploidy, L is the length of the genome, b is the mean number of births per timestep, and T is the number of timesteps. So as to avoid generating such a large amount of random numbers throughout a model run, when a Geonomics model is first created, Geonomics generates and saves (as binary arrays) a large collection of recombination paths (the exact number of which can be parameterized by the user). (A recombination path is a locuswise map along a pair of homologous genomes, represented as 0 and 1, where a switch between 0 and 1 at each locus is drawn as a Bernoulli random variable with success equal to the rate of recombination between that locus and the next.) Then, the paths are repeatedly shuffled and drawn through (think of a deck of cards) and used to index new gametes out of the genomes of reproducing individuals. In this way, free and potentially heterogeneous recombination is approximated. We feel that this is a reasonable approach that balances the needs of the system being modeled with the constraints of the interpreted language being used for the model, and that poses no risk of creating systematic artefacts. However, the user should be aware of two concerns: 1. The total number of recombination paths generated (which is controlled from the parameters file) determines the precision with which truly free recombination at the stipulated genome-wide recombination rates is approximated. 2. The total number of paths also dictates the precision at which close recombinations rates can be differentiated, because all loci with non-zero recombination rates are forced to have at least a single recombination rate among the set of recombination paths that are generated (e.g., the smallest recombination rate that can be modeled when 10000 recombination paths are used is $1/10000 = 10 \exp^{-4}$, such that a locus with a nominal recombination rate of $10 \exp^{-5}$ would be inaccurately modeled as recombining with a probability of $10 \exp^{-4}$). When necessary, Geonomics raises relevant warnings about this at the time that a model is created.

The other approximation of which users should be aware is that of movement across a heterogeneous resistance layer. This is implemented by the movement surface and dispersal surface objects (both of which are identical in structure and function). Conceptually, a movement or dispersal surface is an xx array of VonMises (i.e circular) probability distributions. Each distribution pertains to its corresponding landscape cell, and for an individual located on that cell and about to undergo movement or dispersal it is used as the

distribution from which the individual's direction is drawn. The distributions can be made either uni- or multimodal (i.e. either single VonMises distributions or VonMises mixture distributions), and they represent the probability of an individual in a given cell moving in all directions, based on the environmental values in the surrounding cells. Programmatically, this is implemented using an array of vectors of draws from the distributions, such that the vectors serve as approxiations of the true distributions (much like unbinned histograms). Thus, a movement or dispersal surface is structured as an `xxyxz` numpy array, where x and y are the x- and y-dimensions of the landscape and z is the length of the vectors approximating the true distributions. This structure is built when a model is first created, and then allows us to simulate realistic movement across a complex landscape without the need to run neighborhood-querying functions (which are notoriously costly) during each model run. However, the accuracy of the approximations of the true distributions is a function of the length of the approximation vectors. For complex landscapes, users may wish to ensure the accuracy with which their movement or dispersal surfaces reflect the true heterogeneity in their resistance layers. Geonomics provides a number of built-in visualization functions to aid in heuristic checks of this sort.

9 Getting started

For those interested in using Geonomics, the simplest way is to install via pip (a Python package installer), by calling `$ pip install geonomics`. (Note that Geonomics only uses common, well established Python packages as required dependencies: Numpy [18], Matplotlib [10], Pandas [15], Shapely [5], Scipy [12], Scikit-learn [19], statsmodels [25], PyVCF [1], and Bitarray [24]. The source code is also publicly available on GitHub ([URL_HERE](#)), where it is actively maintained and developed.

The documentation is comprehensive (and will continue to expand as new worked examples are offered and new functionality is added to the package), and is available online at http://htmlpreview.github.io/?https://github.com/drewhart/geonomics_docs/blob/master/built/doc.html. The simplest and advanced use cases are explained above (see section *Design, structure, and function*), as well as in the documentation. The documentation provides a detailed section explaining the meaning, default values, and usage for every parameter in a Geonomics parameter file. Details about the usage of each function are available both in the documentation and in the docstrings (the latter of which can be accessed by calling Python's `help()` function on a function of interest).

10 Conclusion

- very very brief summary of, once again, Geonomics' design-purposed ability - think it should be very helpful for running simulations of high utility to molecular ecology, conservation, etc. - should be valuable for theoretical, methodological, empirical, and applied research - expandibility planned in advance (e.g. multi-species models foreseen)

11 Works Cited

References

- [1] PyVCF - A Variant Call Format Parser for Python — PyVCF 0.6.8 documentation.
- [2] Rowan D. H. Barrett, Stefan Laurent, Ricardo Mallarino, Susanne P. Pfeifer, Charles C. Y. Xu, Matthieu Foll, Kazumasa Wakamatsu, Jonathan S. Duke-Cohan, Jeffrey D. Jensen, and Hopi E. Hoekstra. Linking a mutation to survival in wild mice. *Science*, 363(6426):499–504, February 2019.
- [3] Rachael A. Bay, Ryan J. Harrigan, Vinh Le Underwood, H. Lisle Gibbs, Thomas B. Smith, and Kristen Ruegg. Genomic signals of selection predict climate-driven population declines in a migratory bird. *Science*, 359(6371):83–86, January 2018.

- [4] R. A. Fisher. XXI.—On the Dominance Ratio. *Proceedings of the Royal Society of Edinburgh*, 42:321–341, 1923.
- [5] Sean Gillies et al. Shapely: manipulation and analysis of geometric objects, 2007–.
- [6] F. Guillaume and J. Rougemont. Nemo: an evolutionary and population genetics programming framework. *Bioinformatics*, 22(20):2556–2557, October 2006.
- [7] Benjamin C Haller and Philipp W Messer. SLiM: An Evolutionary Simulation Framework. page 517.
- [8] Benjamin C. Haller and Philipp W. Messer. SLiM 2: Flexible, Interactive Forward Genetic Simulations. *Molecular Biology and Evolution*, 34(1):230–240, January 2017.
- [9] Daniel L. Hartl and Andrew G. Clark. *Principles of Population Genetics: Fourth Edition*. Sinauer Associates, Inc. Publishers, Sunderland, Massachusetts, fourth edition, 2007.
- [10] J. D. Hunter. Matplotlib: A 2d Graphics Environment. *Computing in Science Engineering*, 9(3):90–95, May 2007.
- [11] National Cancer Institute. Genetic Simulation Resources Home.
- [12] Eric Jones, Travis Oliphant, Pearu Peterson, et al. SciPy: Open source scientific tools for Python. [Online; accessed 1 today].
- [13] Erin L. Landguth, Andrew Bearlin, Casey C. Day, and Jason Dunham. CDMetaPOP: an individual-based, eco-evolutionary model for spatially explicit simulation of landscape demogenetics. *Methods in Ecology and Evolution*, 8(1):4–11, January 2017.
- [14] Erin L. Landguth and S. A. Cushman. cdpop: A spatially explicit cost distance population genetics program: COMPUTER PROGRAM NOTE. *Molecular Ecology Resources*, 10(1):156–161, January 2010.
- [15] Wes McKinney. Data structures for statistical computing in python. In Stéfan van der Walt and Jarrod Millman, editors, *Proceedings of the 9th Python in Science Conference*, pages 51 – 56, 2010.
- [16] Philipp W. Messer. SLiM: Simulating Evolution with Selection and Linkage. *Genetics*, 194(4):1037–1039, August 2013.
- [17] S. Neuenschwander, F. Hospital, F. Guillaume, and J. Goudet. quantiNemo: an individual-based program to simulate quantitative traits with explicit genetic architecture in a dynamic metapopulation. *Bioinformatics*, 24(13):1552–1553, July 2008.
- [18] Travis E. Oliphant. *Guide to NumPy*. CreateSpace Independent Publishing Platform, USA, 2nd edition, 2015.
- [19] Fabian Pedregosa, Gael Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, Jake Vanderplas, Alexandre Passos, and David Cournapeau. Scikit-learn: Machine Learning in Python. *MACHINE LEARNING IN PYTHON*, page 6.
- [20] Fanie Pelletier. Testing evolutionary predictions in wild mice. *Science*, 363(6426):452–453, February 2019.
- [21] B. Peng and M. Kimmel. simuPOP: a forward-time population genetics simulation environment. *Bioinformatics*, 21(18):3686–3687, September 2005.
- [22] Adam Porter. ClineFit v. 2.0- β 1-beta version, User's Manual. August 2013.

- [23] François Rebaudo, Arnaud Le Rouzic, Stéphane Dupas, Jean-François Silvain, Myriam Harry, and Olivier Dangles. SimAdapt: an individual-based genetic model for simulating landscape management impacts on populations. *Methods in Ecology and Evolution*, 4(6):595–600, June 2013.
- [24] Ilan Schnell. efficient arrays of booleans for Python. Contribute to ilanschnell/bitarray development by creating an account on GitHub, February 2019. original-date: 2011-01-20T03:20:11Z.
- [25] Skipper Seabold and Josef Perktold. Statsmodels: Econometric and statistical modeling with python. In *9th Python in Science Conference*, 2010.
- [26] Sewall Wright. Evolution in Mendelian populations. *Genetics*, 16(2):97–159, January 1930.

12 Figures and Captions

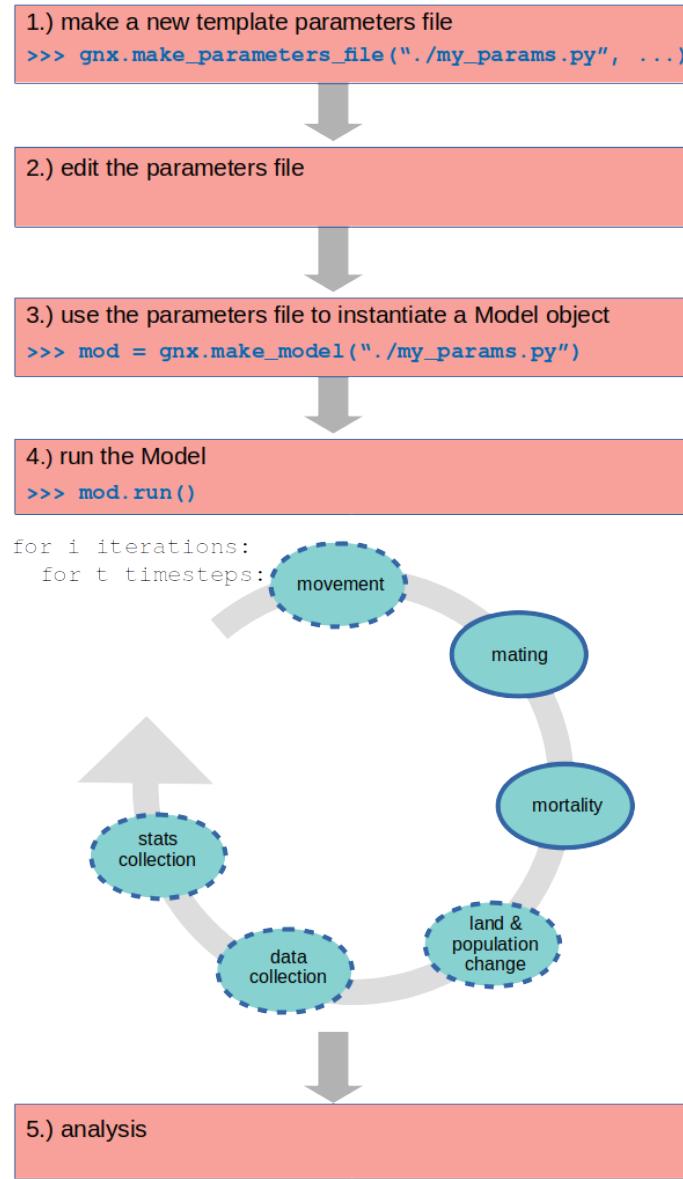


Figure 1: Geonomics flow diagram

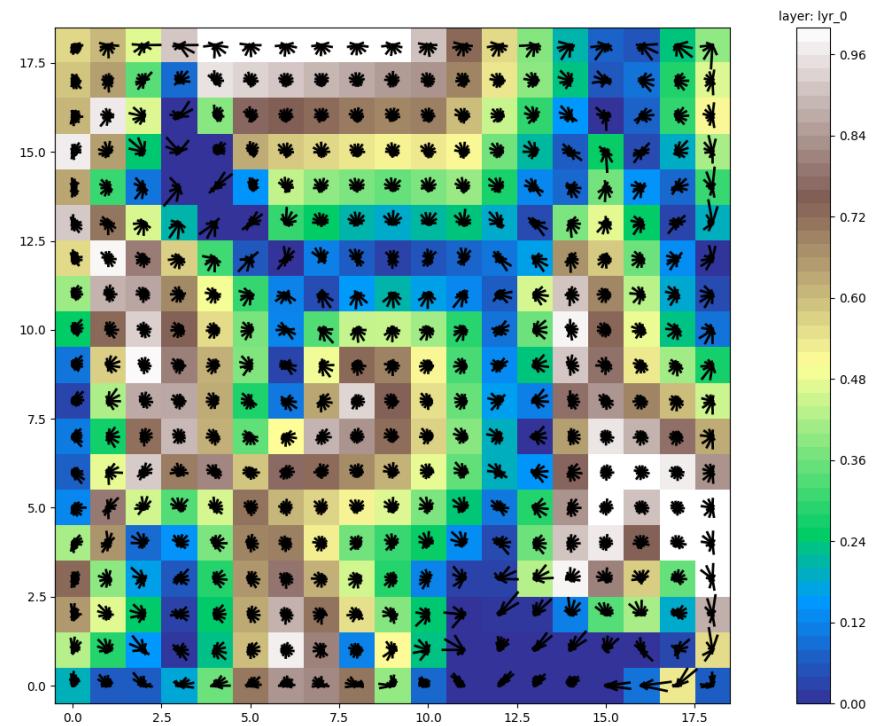


Figure 2: Histograms of the VonMises mixture-distribution approximations in a movement surface, plotted over a small landscape layer on which the surface is based. The taller the bar in a histogram, the higher the probability that an individual located on that histogram's cell will move in a direction within the angular neighborhood of that bar.

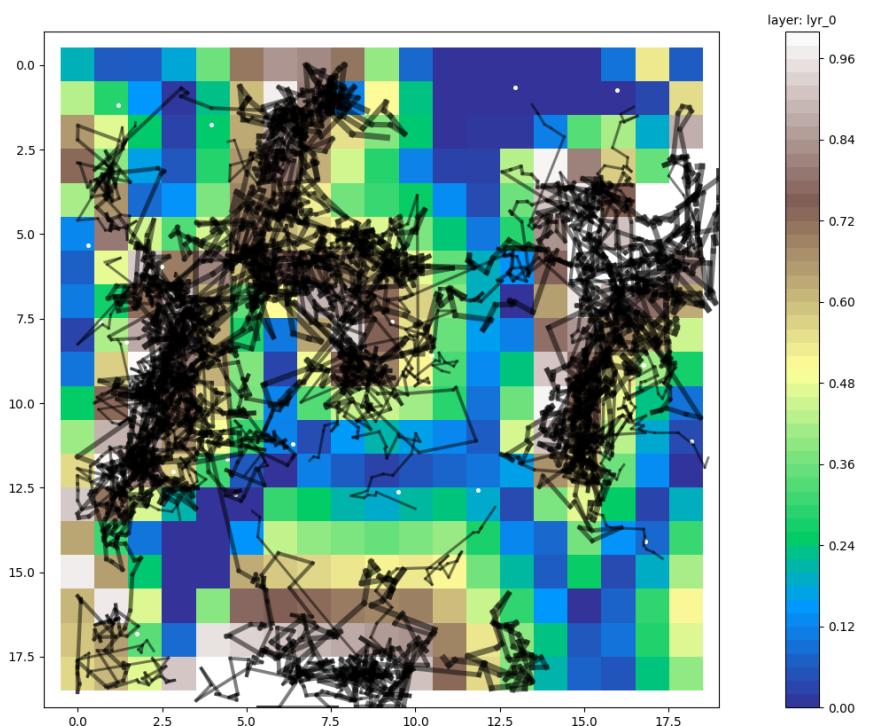


Figure 3: Examples of movement tracks for 25 individuals randomly selected from a species, moving across the movement surface depicted in Figure 2. Each track is 150 steps long, beings at the individuals' starting locations (white dots), and thickens with each increasing timestep. Individudals' preference for higher-suitability regions of the landscape (values with environmental values nearer 1) is evident. Occasional migrations between isolated portions of the landscape can also be seen.

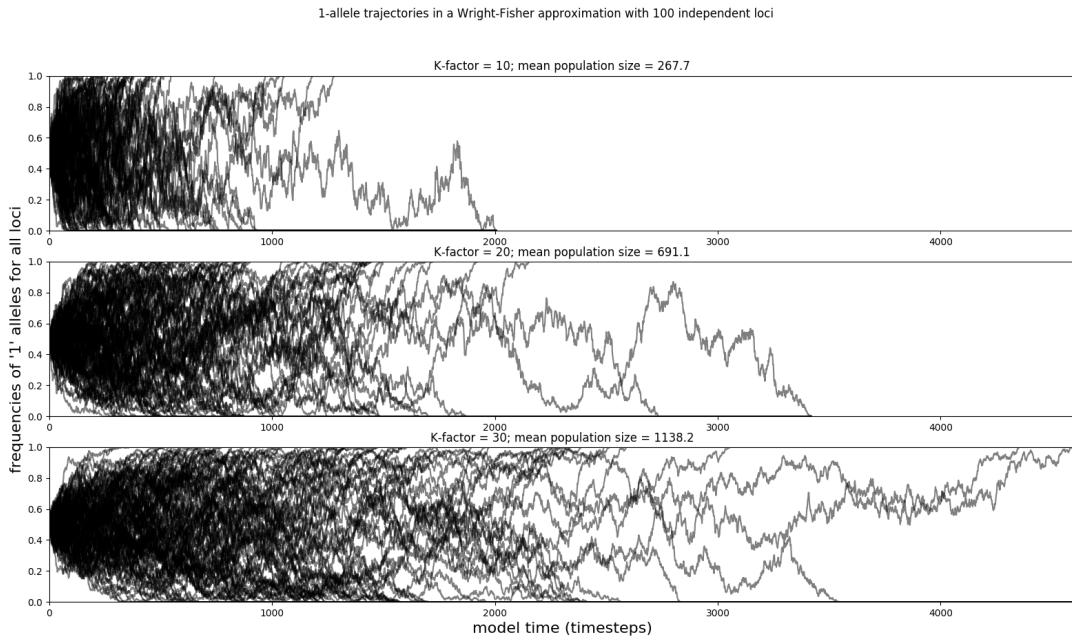


Figure 4: Wright-Fisher test: allele-frequency trajectories; each line shows the trajectory of the '1'-allele for a different locus

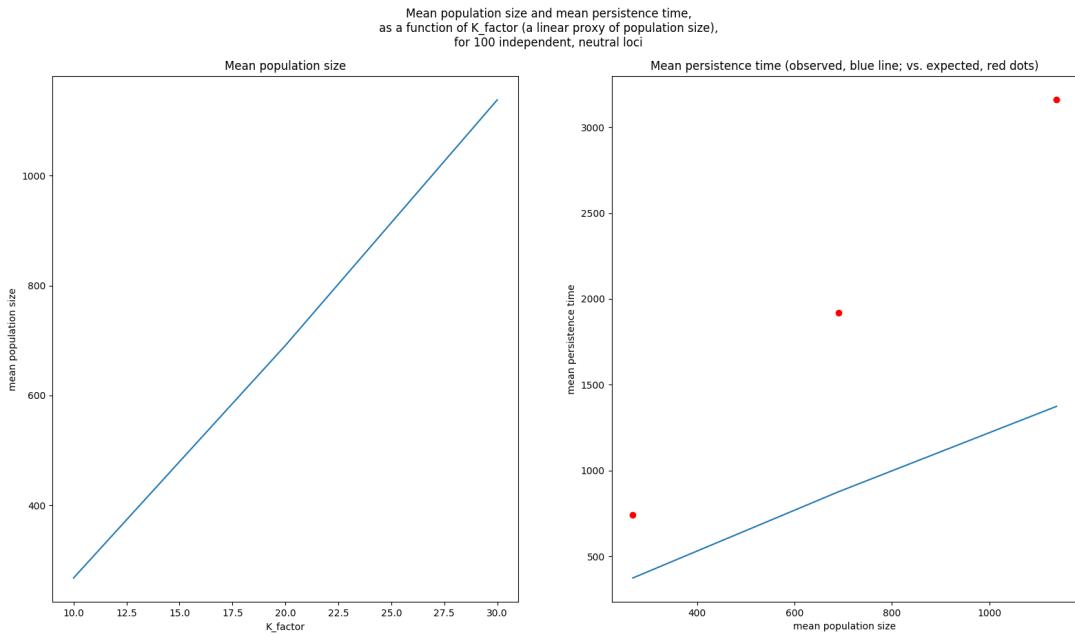


Figure 5: Wright-Fisher test: left: mean population size as a function of 'K_factor' (multiplicative factor determining a species' spatialized carrying capacity and thus equilibrium population size; right: mean persistence time for segregating sites as a function of mean population size)

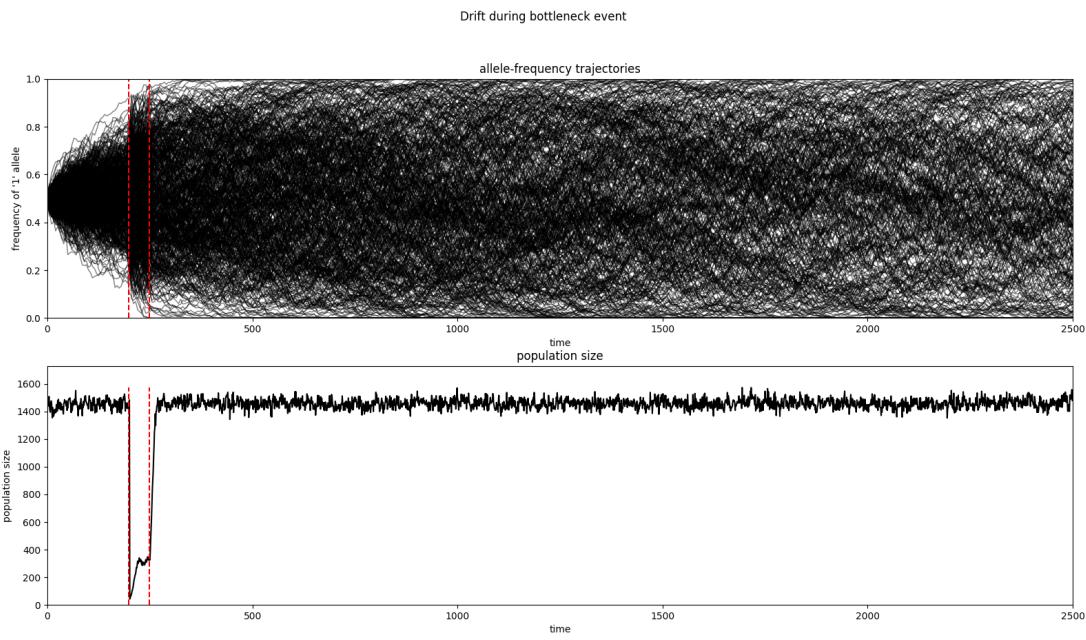


Figure 6: Bottleneck test: 1-allele frequencies (top) and population size (bottom) as a function of time

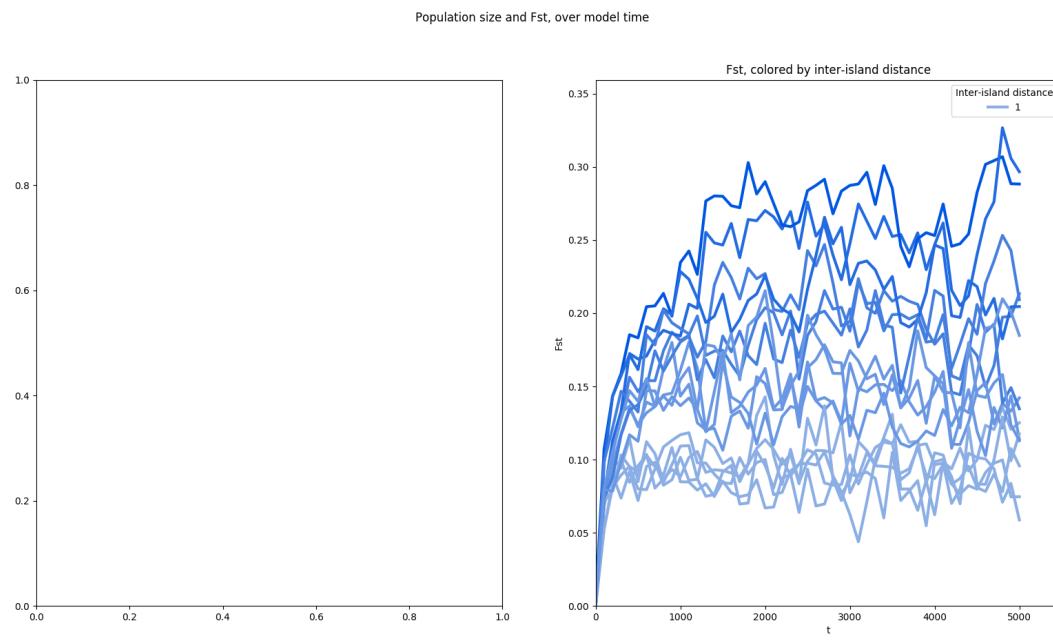


Figure 7: Stepping-stone test: F_{ST} as a function of model time, across increasing inter-island distances (gradually darker shades of blue)

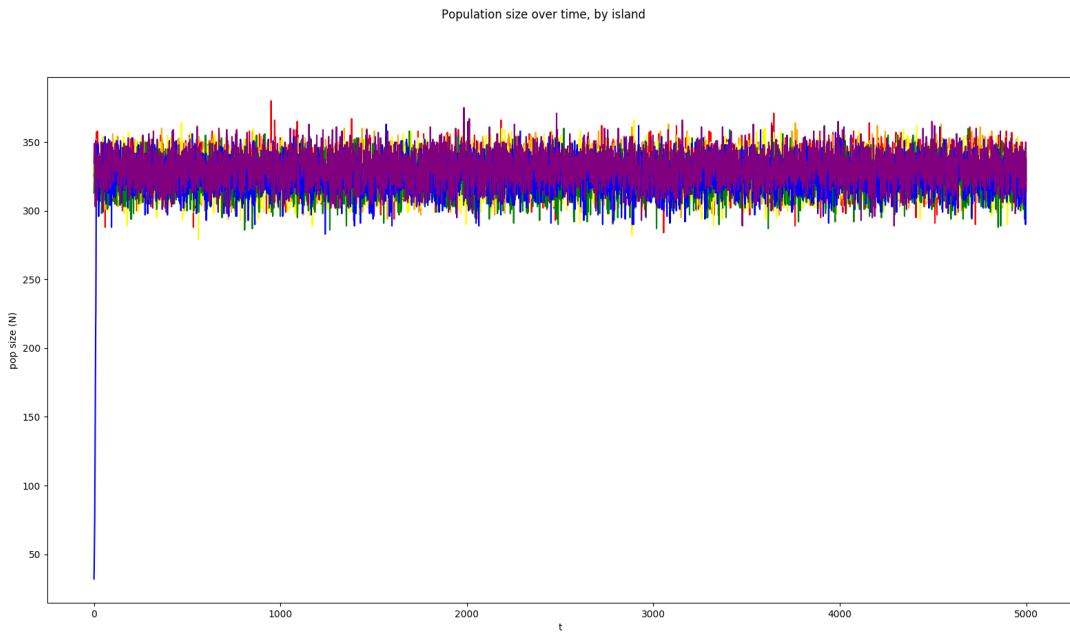


Figure 8: Stepping-stone test: Population size as a function of time, for all 6 islands' populations

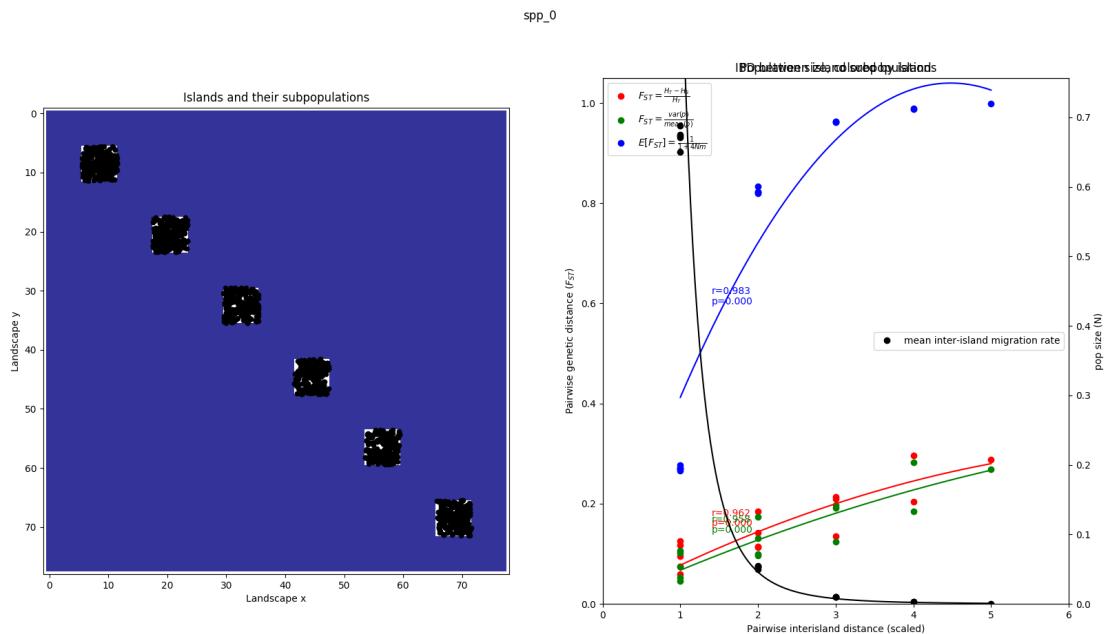


Figure 9: Stepping-stone test: left: Map of all 6 islands' populations at the end of the simulation; right: pairwise F_{ST} (left y-axis; calculated by 2 different formulae) and inter-island migration rate (right y-axis) as a function of inter-island distance; R^2 values and p-values result from quadratic regressions of F_{ST} values on inter-island distance and log-log regression of mean migration rate on inter-island distance

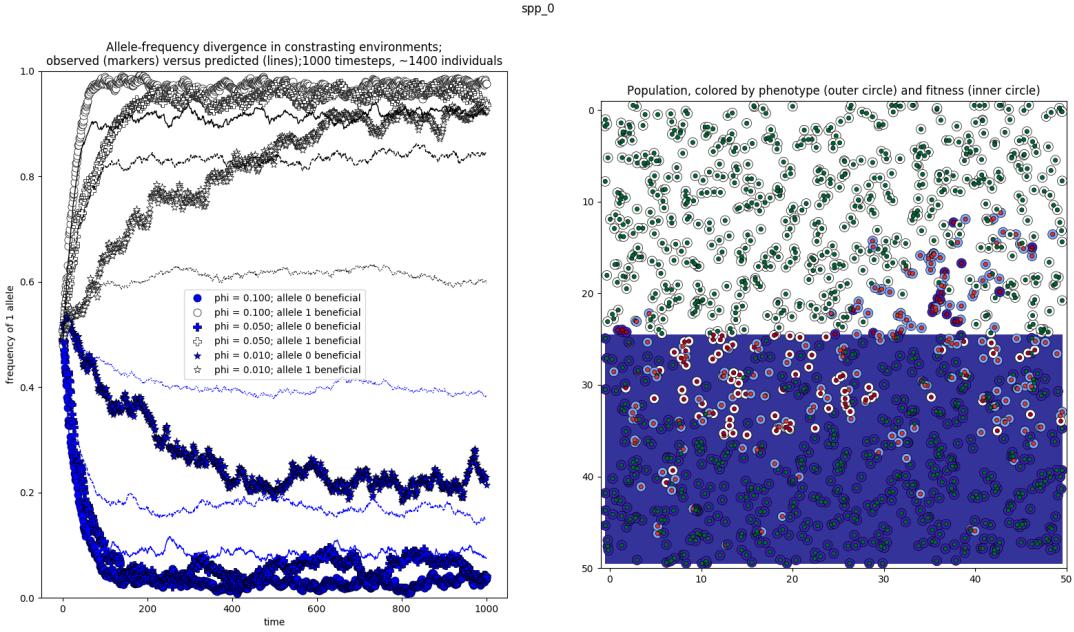


Figure 10: Contrasting-habitat test: left: Observed (markers) versus predicted (lines) allele-frequency trajectories for two contrasting habitats (blue = 0.0-valued; white = 1.0-valued), across three selection coefficients ($\phi = 0.01$: stars; 0.05: crosses; 0.10: circles); right: map of the population after spatially divergent selection at $\phi = 0.10$, with individuals, colored by phenotype (outer circles) and fitness (inner circles), plotted on top of the selective landscape layer (horizontally divided into white and blue halves)

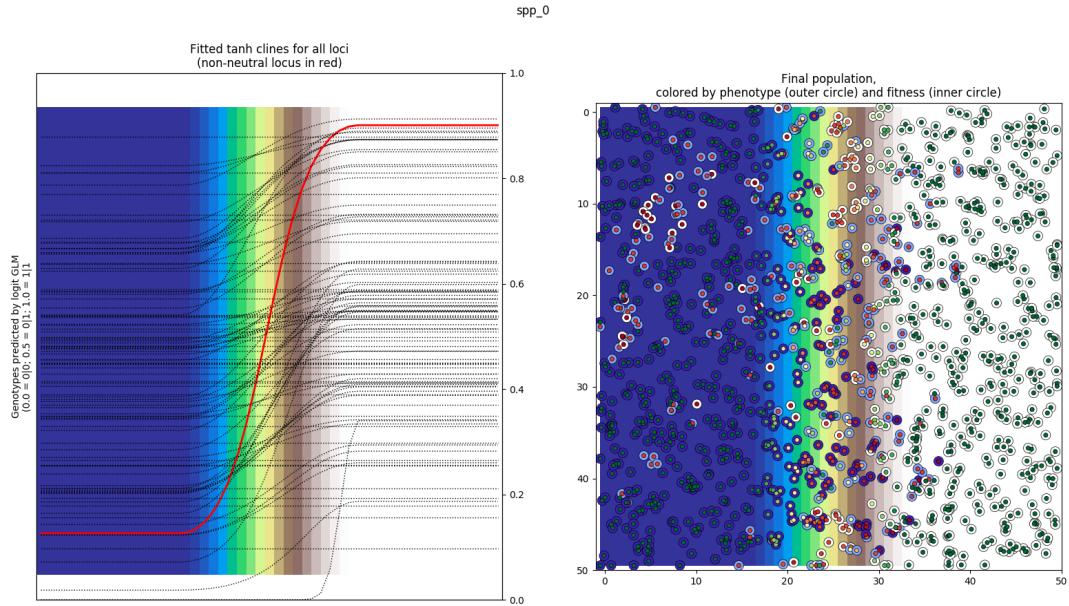


Figure 11: Cline test: left: plot of allele-frequency clines (neutral loci in black, selective locus in bold red) against the selective landscape layer (horizontal gradient from blue to white); right: map of the final population on top of the selective landscape layer, with individuals colored by phenotype (out circles) and fitness (inner circles)

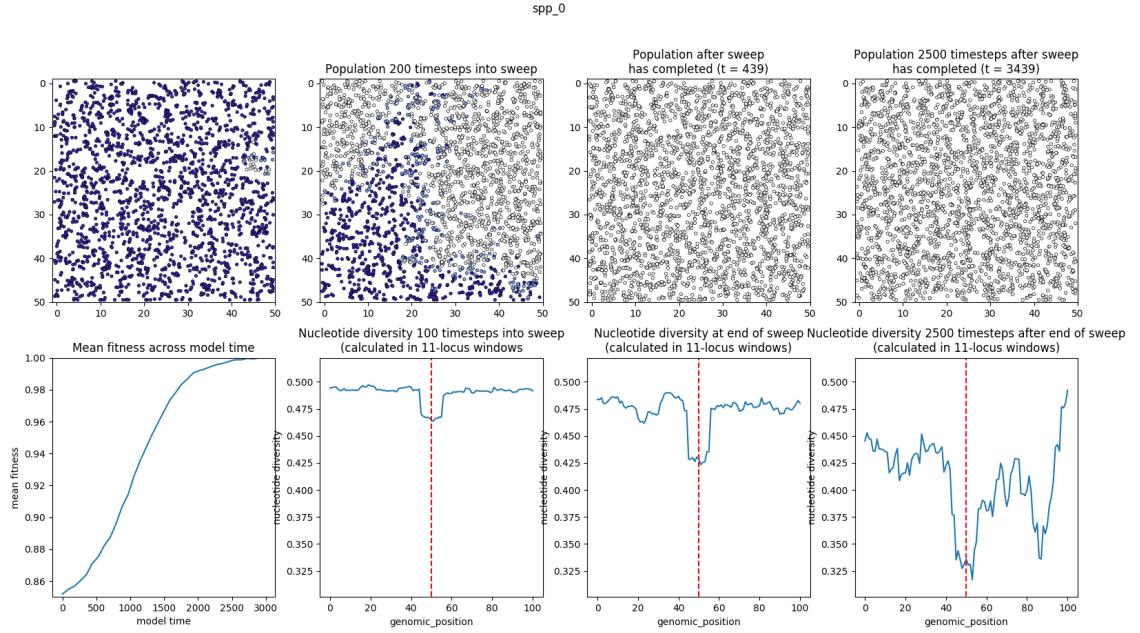


Figure 12: Selective sweep test: top row: Maps of population (colored by phenotype) at various points in model time (from left to right: timestep 0, timestep 200, after completion of sweep, and 2500 timesteps after completion of sweep); bottom row: mean fitness as a function of model time (first plot on left) and genome-wide nucleotide diversity at timestep 200, immediately after completion of sweep, and 2500 timesteps after completion of sweep (second, third, and fourth plots from left)

Neutral genomic evolution across complex landscape with _MovementSurface,
(for a ~3291-individual species with 100 loci)

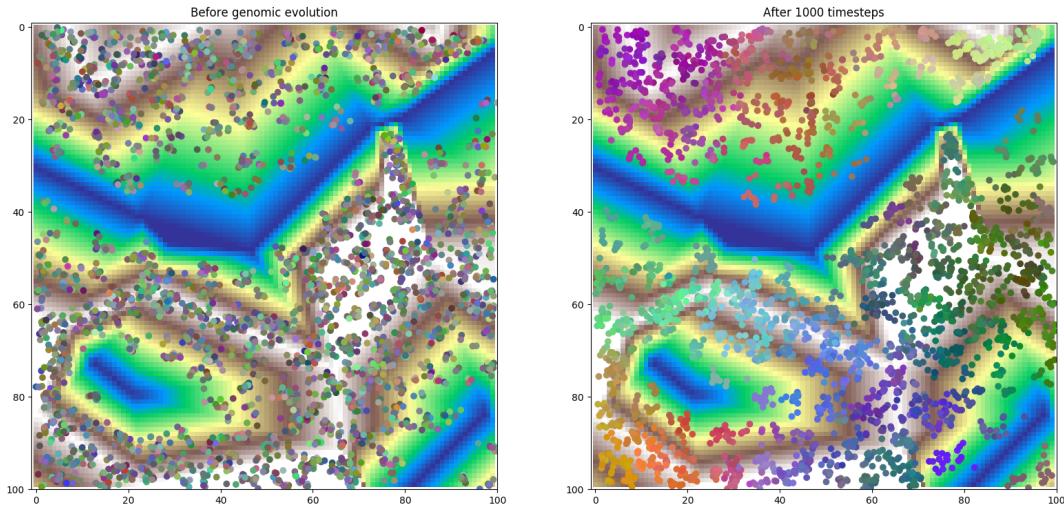


Figure 13: Spatial structure in a species evolving on a complex landscape layer serving as a movement surface, before (left) and after (right) 1000 timesteps of neutral evolution. Individuals' colors are derived from their values for the first three PCs of a genetic PCA (each PC scaled to $0 \leq value \leq$, then used to assign RGB values)

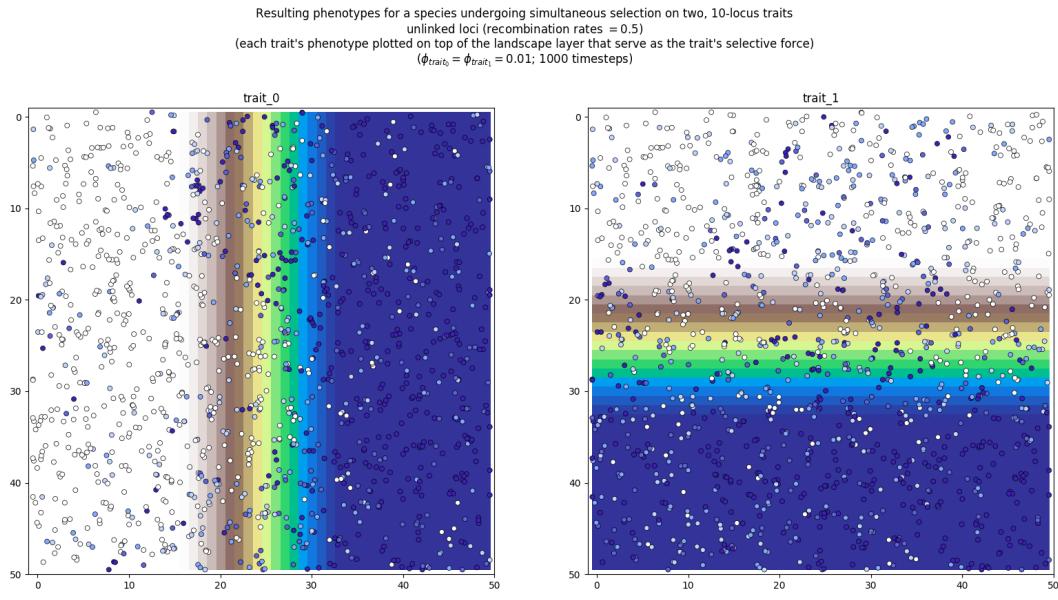


Figure 14: Results of simultaneous selection on two traits with divergent maps of selective force. Each trait has 10 unlinked loci and a selection coefficient of $\phi = 0.05$

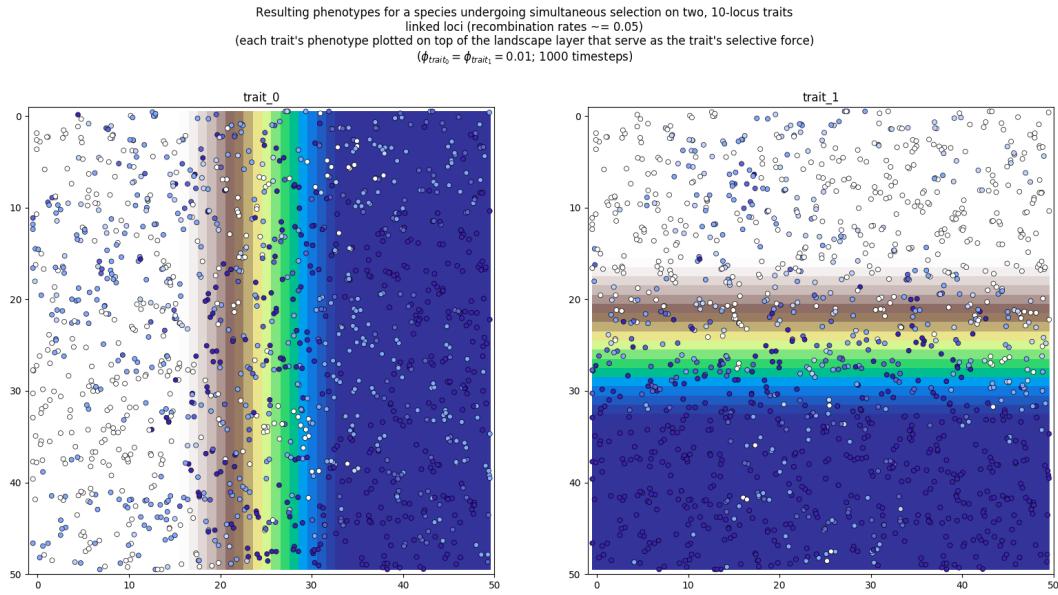


Figure 15: Results of simultaneous selection on two traits with divergent maps of selective force. Each trait has 10 linked loci (recombination rate = 0.05) and a selection coefficient of $\phi = 0.05$

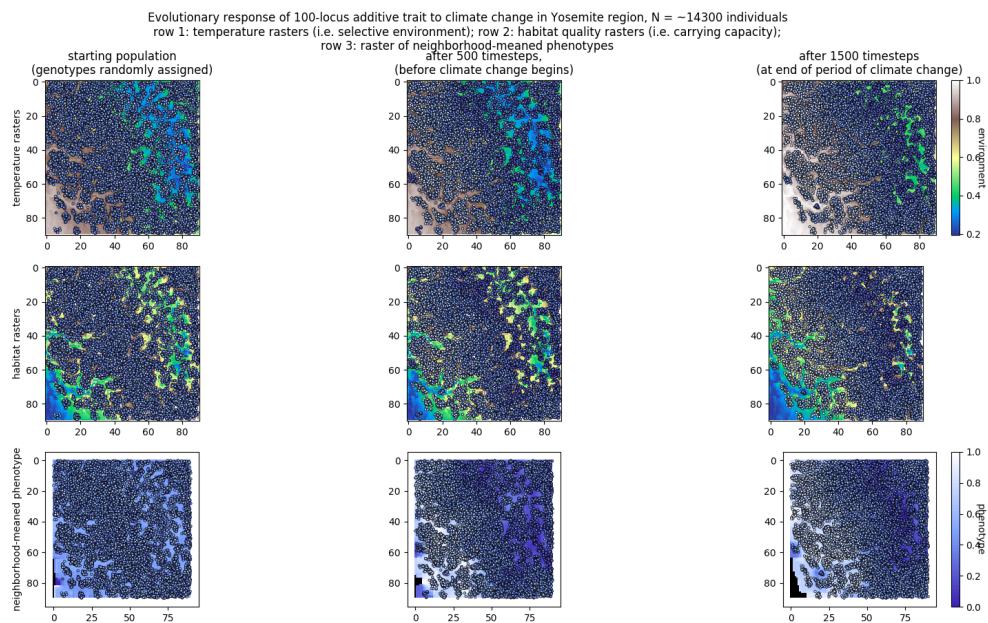


Figure 16: Temperature rasters (top row), habitat rasters (middle row), rasters of neighborhood-meaned phenotype (bottom row) at timesteps 0 (left column), 500 (before beginning of climate change; center column), and 1500 (after climate change; right column) for a species with a 100-gene trait adapted to temperature